

ANA670 Applied Optimization

Assignment Module Four

First Name	Randall
Last Name	Crawford
ID#	901012194
Email Address	rcc_0149@yahoo.com

Table of Contents

How to submit your Assignment	1
P1 - Genetic Algorithms [30 points]	2
P2 - PSO [65 points]	4
P3 – Text classification (30 points)	6

How to submit your Assignment

After filling all the parts in this file, please follow the following steps.

- 1) Add your name and ID to the first page.
- 2) Enter your answers and save the file as a PDF (pdf)
- 3) Rename the file as
YOUR First Name - YOUR Last Name - ID - ANA670 – HW4.pdf

Example: John - Smith - ID - ANA670 – HW4.pdf

- 4) Upload the file and submit it (only using BrightSpace)

P1 - Genetic Algorithms [30 points]

The solution space for a Genetic Algorithm is encoded in terms the algorithm can work with.

P1.Q1: What are the inputs and outputs of this encoding process? (Cite an example from academic literature, e.g., a paper you found on [scholar.google.com](#) or [arxiv.org](#))

P1.Q2: What essential operations does a Genetic Algorithm perform on the encoded information (explain them a little in your own words)?

P1.Q3: Is it necessary to decode the final solution (why or why not)?

Explain your reasoning in more than 75 words but less than 150.

(Write your answer yourself. AI generated text in the box below will usually be returned as “does not make sense”)

Your Answer

P1.Q1. I would like to consider a modern application of GAs in the optimization of neural network architectures, a field where Holland's binary encoding has been adapted and enhanced. This example is drawn from contemporary research that builds on his principles, such as work seen in neuroevolutionary studies (e.g., approaches by researchers like Kenneth O. Stanley and Risto Miikkulainen with Evolving Neural Networks through Augmenting Topologies on scholar.google.com). The method is referred to as Neuro-Evolution of Augmenting Topologies (NEAT), which modernizes Holland's ideas.

Regarding inputs for the encoding process:

Problem Definition: The task is to optimize the architecture of a neural network (e.g., number of layers, neurons per layer) for a specific machine learning task, like image classification.

Solution Parameters: Variables include the presence or absence of connections, node types (e.g., input, hidden, output), and layer configurations.

Encoding Scheme: Building on Holland's binary encoding, the architecture is represented as a binary string where each bit indicates the presence (1) or absence (0) of a connection or node feature, expanded with additional bits to encode more complex structures (e.g., a 50-bit string where segments represent different network components).

Regarding outputs for the encoding process:

Encoded Individuals: A population of binary strings (e.g., “101100110010...”), each representing a unique neural network configuration.

Fitness Function: The accuracy of the neural network on a validation dataset (e.g., MNIST digit recognition) serves as the fitness metric, evaluated through training and testing.

This modernizes Holland's binary encoding by applying it to a dynamic, complex domain like neural networks, where his schema theorem (about building blocks) guides the search for effective architecture.

P1.Q2. In this neuroevolutionary context, the GA performs operations that evolve the binary-encoded network architectures:

Selection: The algorithm chooses the best-performing networks (those with higher classification accuracy) to reproduce, like picking the smartest designs to carry forward.

Crossover: Two parent binary strings swap segments (e.g., exchanging parts of “10110011” and “01001100” to create new combinations like “10101100”), blending network structures to explore new configurations.

Mutation: Random bit flips (e.g., changing “0” to “1” in “10110011” to get “10110001”) introduce variations, allowing the algorithm to discover novel connections or nodes.

Evaluation: Each binary string is decoded into a network, trained on data, and scored based on its performance, guiding the evolutionary process.

In my own words, it's like breeding neural designs where the best ones are selected, mixed up a bit, given random tweaks, and then tested on a task. Repeating the process until a top-notch network is achieved.

P1.Q3. In this neural network optimization scenario, yes, decoding is necessary. The final solution is a binary string (e.g., “101100110010...”) representing the network architecture. To implement it, this string must be decoded into a functional neural network by mapping bits to connections, nodes, and layers, then training it with weights. The GA evolves the architecture in an encoded form for efficiency, but the practical use (e.g., deploying the network for image classification) requires a decoded structure. Without decoding, the binary string is just a code—meaningless until turned into a working model.

This aligns with Holland's emphasis on representation, modernized to handle the complexity of neural architecture while still requiring a translation step for real-world application.

P2 - PSO [65 points]

Have a look at the book chapter attached in a PDF to this assignment and answer these questions in your own words:

1. In the chapter, which line of code in **Algorithm 16.1** do you think **Figure 16.1** describes (explain why)?
2. How does the following in **Figure 16.1** help the algorithm arrive at a solution?
 - (I) “social velocity”
 - (II) “inertia velocity”
 - (III) “cognitive velocity”
 - (IV) “new velocity”
3. What is happening during the transition from time t to time $t+1$ in **Figure 16.1**?
4. What is the problem described in section **16.5.1**, and how is it resolved?

Your Summary

- 1) Line of Code: “update the position using equation (16.1);”
 - Equation: $x_i(t + 1) = x_i(t) + v_i(t + 1)$
 - Figure 16.1 visually demonstrates the effect of the velocity update on a particle’s position, where the new position $x_i(t + 1)$ is calculated by adding the updated velocity $v_i(t + 1)$ to the current position $x_i(t)$.
- 2) I. The social velocity represents the influence of the global best position $\hat{y}(t)$ found by the entire swarm. It helps the algorithm arrive at a solution by guiding each particle toward the best-known position in the search space, encouraging convergence toward a potentially optimal solution. This collective knowledge accelerates the swarm’s movement toward promising regions.

II. The inertia velocity acts as a memory of the particle’s prior flight direction, carrying forward momentum from the last step. This helps maintain a consistent movement pattern unless significantly altered by the cognitive or social components.

III. The cognitive velocity aids the algorithm by allowing each particle to learn from its personal experience, pulling it back toward positions where it performed well. This individual memory helps prevent the particle from straying too far from effective areas and contributes to refining the solution by balancing exploration with the particle’s past successes.
- IV. The new velocity drives the position update and is the net effect of the three aforementioned components. It helps the algorithm arrive at a solution by determining the direction and magnitude of each particle’s movement, combining individual and social influences to iteratively steer the swarm toward an optimal region.

- 3) During the transition from time t to time $t + 1$ in Figure 16.1, the particle's position is updated based on its velocity, which is influenced by its previous movement, personal best position, and the global best position.

Initial State (Time t): The figure shows the particle's current position $x(t)$ at time t , along with its personal best position $y(t)$ and the global best position $\hat{y}(t)$. The velocity $v(t)$ reflects the particle's previous direction of movement.

Velocity Update: The velocity is recalculated using equation (16.2) for gbest PSO, incorporating three components:

- The inertia term $v_{ij}(t)$, which maintains the particle's previous direction (momentum).
- The cognitive component $c_1 r_{1j}(t)[y_{ij}(t) - x_{ij}(t)]$, which pulls the particle toward its own best position.
- The social component $c_2 r_{2j}(t)[\hat{y}_j(t) - x_{ij}(t)]$, which attracts the particle toward the swarm's best position. The updated velocity $v(t + 1)$ is a stochastic combination of these factors, adjusted by random values r_{1j} and r_{2j} .

Position Update: The new position $x(t + 1)$ is computed using equation (16.1): $x_i(t + 1) = x_i(t) + v_i(t + 1)$. Figure 16.1 illustrates this as the particle moves closer to the global best $\hat{y}(t)$, influenced by the combined effect of the velocity components.

4) Guaranteed Convergence Particle Swarm Optimization (GCPSO)

Problem: The standard PSO, particularly the gbest version, can suffer from premature convergence or stagnation, where the swarm gets trapped in local optima or fails to refine solutions effectively. This issue arises because the velocity update equation (16.2) relies on a balance of cognitive and social components, but the algorithm lacks a guaranteed mechanism to ensure convergence to a local optimum, especially in complex, multimodal search spaces.

Resolution: The GCPSO, introduced as a variation, modifies the velocity update to guarantee convergence to a local optimum. Instead of using the standard velocity equation (16.2), GCPSO adjusts the particle's movement by directly updating its position toward a weighted combination of its personal best y_i and the global best \hat{y} , with a constriction factor to control the step size. The key change is in the velocity update, where the particle's position is iteratively refined to ensure it approaches a stable point. This is achieved by incorporating a mechanism (often a constriction coefficient) that dampens the velocity, preventing divergence and ensuring that particles converge to a local optimum.

P3 – Text classification (30 points)

- (I) Describe the approach of [this paper](#) in applying PSO to Document classification.
- (II) Why do they combine PSO with k-means?
- (III) To the best of your understanding what would a solution generated by the algorithm look like? (The paper offers some information about this but Table 1 only gives some meta-results). Also, what is its meaning of such a solution?
- (IV) Can you think of a possible application or use case for this (e.g. an application that a startup business could use profitably)?
- (V) What do you think about using clustering algorithms on text classification? (Does it surprise you? Why or why not?)

Your Summary

- (I) The paper presents a semantic document clustering approach that integrates Particle Swarm Optimization (PSO) with the Universal Networking Language (UNL) to enhance document classification. The process begins with preprocessing of documents, which includes tokenization, stop word removal, and stemming to prepare the text for representation. Unlike traditional methods that rely on term frequency (TF) or TF-IDF, the paper introduces a novel representation using UNL, which captures semantic relationships between words through a semantic graph. Each document is represented as a feature vector where Universal Words (UWs) serve as nodes, and the number of links (semantic relations) between them determines the weight of each UW.

The classification (or clustering) is performed by using a hybrid PSO + K-means algorithm. Initially, PSO is employed to globally search for optimal cluster centroids by treating each particle as a potential solution representing cluster centroids. The fitness of each particle is evaluated based on the intra-cluster similarity (minimizing distances within clusters) and inter-cluster separation (maximizing distances between clusters). The PSO algorithm iteratively updates particle positions and velocities using equations that balance local and global search capabilities, guided by the best positions found by individual particles and the swarm. Once PSO converges or reaches a maximum iteration, the resulting centroids are used as initial seeds for the K-means algorithm, which refines the clustering with faster local convergence. This hybrid approach leverages PSO's global optimization with K-means' efficiency, leading to improved clustering accuracy, as demonstrated by the experimental results where the UNL-based method achieved a 97.63% accuracy compared to 64.28% and 85.71% for TF and TF-IDF methods, respectively.
- (II) The combination of PSO with K-means is motivated by the complementary strengths of both algorithms. PSO is a population-based stochastic optimization technique that excels at global search, exploring a wide range of potential solutions to avoid getting trapped in local optima. However, it can be computationally intensive and may require more iterations to converge precisely. On the other hand, K-means is a fast-converging algorithm that performs well in local optimization but is sensitive to initial centroid placement and can converge to suboptimal solutions if the initial seeds are poor.

By integrating the two, the hybrid algorithm first uses PSO to perform a global search, identifying promising cluster centroids that are close to the optimal solution. These centroids are then fed into the K-means algorithm, which refines the clustering with its efficient local optimization. This synergy allows the method to benefit from PSO's ability to explore the solution space broadly while utilizing K-means' speed to fine-tune the results, ultimately improving both the quality and efficiency of document clustering. The paper highlights that this hybrid approach outperforms standalone PSO or K-means, as evidenced by the superior accuracy and cluster quality metrics in the experiments.

- (III) Based on the paper, a solution generated by the hybrid PSO + K-means algorithm would manifest as a set of clusters, where each cluster groups documents with similar semantic content. The solution is represented by the final cluster centroids (points in the multidimensional feature vector space) determined after the PSO phase identifies approximate optimal centroids and K-means refines them. Each document is assigned to the cluster whose centroid it is closest to, based on a distance metric like cosine similarity. The feature vectors, constructed using UNL, reflect semantic relationships, so the clusters are not just based on word frequency but on conceptual similarity.

Table 1 provides meta-results, showing metrics like intra-cluster similarity (e.g., 0.7442 for UNL), inter-cluster similarity (e.g., 0.7340 for UNL), and accuracy (e.g., 97.63% for UNL). These indicate a solution where clusters are compact (high intra-cluster similarity), well-separated (low inter-cluster similarity), and highly accurate in correctly classifying documents. Visually, one might imagine a scatter plot of document vectors in a high-dimensional space, with clear, distinct groupings around the centroids, each representing a semantic category. The meaning of such a solution lies in its ability to organize a large document collection into meaningful, semantically coherent groups, facilitating tasks like information retrieval, topic extraction, or automated document management with high precision.

- (IV) This approach has several promising applications, particularly for a startup business. One potential use case is in content management and recommendation systems for a digital media or e-learning platform. A startup could use this semantic clustering to automatically categorize articles, videos, or courses into topics (e.g., machine learning, history, coding) based on their semantic content rather than just keywords. This would enable personal recommendations, improving user engagement and retention, which could translate into increased subscription revenue.

For example, a startup like an online education platform could process a vast library of educational resources, clustering them into coherent curricula or skill tracks. The high accuracy of the UNL-based method (97.63%) ensures that related but differently worded content is grouped together, enhancing the platform's searchability and user experience. Additionally, the system could identify gaps in content coverage, guiding the startup to create new materials profitably. By offering precise, semantically rich content organization, the startup could attract premium users or partnerships with educational institutions, turning the technology into a profitable asset.

- (V) Using clustering algorithms for text classification is an intriguing and logical extension of unsupervised learning techniques, and it doesn't surprise me given the nature of text data and the goals of classification. Clustering is inherently about grouping similar items, and in text classification, the challenge is to assign documents to categories based on their content. Since labeled data for supervised classification might be scarce or expensive to obtain, unsupervised clustering offers a way to discover natural groupings in the data, which can then be manually labeled or used as a starting point for supervised methods.

The paper's use of PSO and K-means with semantic enhancement via UNL aligns with this idea, showing how clustering can be adapted for classification by leveraging semantic relationships to improve grouping quality. This isn't surprising because text data is rich with latent structures (e.g., topics, themes) that clustering can uncover, especially with advanced representations like UNL that address limitations of traditional bag-of-words models (e.g., synonymy, polysemy). What's impressive but not entirely unexpected is the significant accuracy boost (97.63%) over TF-based methods, reflecting how semantic awareness can refine clustering outcomes. The approach makes sense in contexts where predefined categories are unclear or evolving, making it a powerful tool for exploratory analysis or dynamic classification systems.

The End